



Optimization in Machine Learning: Lecture 7

Solving Large-scale Problems I

by Xiaolin Huang

xiaolinhuang@sjtu.edu.cn

SEIEE 2-429

Institute of Image Processing and Pattern Recognition

<http://www.pami.sjtu.edu.cn/>



上海交通大學

SHANGHAI JIAO TONG UNIVERSITY

1

Stochastic Gradient Descent

2

SGD with Momentum

3

Limited-memory BFGS

4

Coordinate Update



1

Stochastic Gradient Descent

2

SGD with Momentum

3

Limited-memory BFGS

4

Coordinate Update



Why stochastic



- starting from least squares

$$\min_x \sum_{i=1}^m (x^\top a - b_i)^2 \triangleq \sum_{i=1}^m f_i(x)$$

- gradient descent:

$$x^{k+1} = x^k - \sum_{i=1}^m \nabla f_i(x^k) = x^k - t \sum_{i=1}^m (x^\top a_i - b_i) a_i$$

- when m is large, it is expensive to iteratively calculate $\nabla f = \sum \nabla f_i$
- stochastically choose a subset from the training set
- use stochastic gradient instead of full gradient



$$\nabla f \approx \nabla f_J = \sum_{i \in J} \nabla f_i$$



Why stochastic



- intuitive impression
 - it sounds reasonable
 - at least, it makes solving large-scale problem feasible
 - performance, convergence, speed?



Convergence analysis



$$\min_x \sum_{i=1}^m (x^\top a_i - b_i)^2$$

- full gradient descent

$$x^{k+1} = x^k - t \sum_{i=1}^m (x^\top a_i - b_i) a_i$$

- stochastic gradient descent: simplest case

$$x^{k+1} = x^k - t \underbrace{(x^\top a_i - b_i) a_i}_{\text{random}}$$



i is randomly selected

$v_i = (x^\top a_i - b_i) a_i$ is a random

Convergence analysis



- notice that $E(v_i) = \nabla f(x)$ (a unbiased estimator)

$$\text{Var}(v_i) = E(\|v\|_2^2) - \|E(v)\|_2^2$$

- with the boundedness condition on the Hessian

$$\nabla^2 f(x) \leq MI, \quad \forall x \in S$$

we have

$$f(x - tv_i) \leq f(x) - t\nabla f^T v_i + \frac{Mt^2}{2} \|v_i\|_2^2$$

- taking expectation on both side

$$\begin{aligned} Ef(x^{k+1}) &\leq f(x^k) - t\nabla f^T E(v_i) + \frac{Mt^2}{2} E(\|v_i\|_2^2) \\ &= f(x^k) - t\|\nabla f(x^k)\|_2^2 + \frac{Mt^2}{2} (\|\nabla f(x^k)\|_2^2 + \text{Var}(v_i)) \end{aligned}$$

Convergence analysis



$$\begin{aligned}
 Ef(x^{k+1}) &\leq f(x^k) - t\|\nabla f(x^k)\|_2^2 + \frac{Mt^2}{2}(\|\nabla f(x^k)\|_2^2 + \text{Var}(v_i)) \\
 &\leq f(x^k) - \frac{t}{2}\|\nabla f(x^k)\|_2^2 + \frac{t}{2}\text{Var}(v_i) \quad (\text{if we choose } Mt \leq 1)
 \end{aligned}$$

SGD is not monotonic

- furthermore, with convexity, $f(y) \geq f(x) + \nabla f(x)^\top(y - x)$

$$\begin{aligned}
 Ef(x^{k+1}) &\leq f(x^*) + \nabla f(x^k)^\top(x^k - x^*) - \frac{t}{2}\|\nabla f(x^k)\|_2^2 + \frac{t}{2}\text{Var}(v_i) \\
 &\leq f(x^*) + E(v_i)^\top(x^k - x^*) - \frac{t}{2}\|E(v_i)\|_2^2 + t\text{Var}(v_i) \\
 &= f(x^*) + E\left(v_i^\top(x^k - x^*) - \frac{t}{2}\|E(v_i)\|_2^2\right) + t\text{Var}(v_i) \\
 &\leq f(x^*) + E\left(\frac{1}{2t}(\|x^k - x^*\|_2^2 - \|x^{k+1} - x^*\|_2^2)\right) + t\text{Var}(v_i)
 \end{aligned}$$

Convergence analysis



$$Ef(x^{k+1}) \leq f(x^*) + E \left(\frac{1}{2t} \left(\|x^k - x^*\|_2^2 - \|x^{k+1} - x^*\|_2^2 \right) \right) + t\text{Var}(v_i)$$

- by summing the inequality

$$\begin{aligned} \sum_{k=0}^K \left(Ef(x^{k+1}) - f(x^*) \right) &\leq \frac{1}{2t} \left(\|x^0 - x^*\|_2^2 - \|x^K - x^*\|_2^2 \right) + Kt\text{Var}(v_i) \\ &\leq \frac{\|x^0 - x^*\|_2^2}{2t} + Kt\text{Var}(v_i) \end{aligned}$$

- for the average result $\bar{x}_K = \sum_{i=0}^K x^k$

$$Ef(\bar{x}_K) \leq f(x^*) + \frac{\|x^0 - x^*\|_2^2}{2tk} + t\text{Var}(v_i)$$

1

Stochastic Gradient Descent

2

SGD with Momentum

3

Limited-memory BFGS

4


Coordinate Update



Revisit gradient descent



- SGD suffers similar problem as full gradient descent



- how to improve?
 - inexact, stochastic, gradient-free,
 -

SGD with momentum



- Nesterov acceleration


$$z_{k+1} = x_k - t_k \nabla f(x_k)$$

$$x_{k+1} = z_{k+1} + \delta_k (z_{k+1} - z_k)$$

- linear combination of the stochastic gradient and the previous update

$$z_{k+1} = x_k - t_k \sum_{i \in J} \nabla f_i(w_k)$$


$$x_{k+1} = z_{k+1} + \delta_k (z_{k+1} - z_k)$$



Heuristic for learning rate





- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem



Heuristic for learning rate



- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem
- learning rate decay


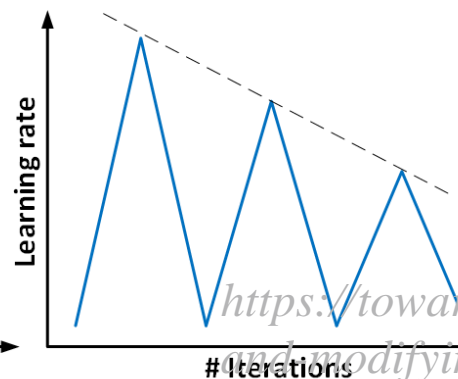



<https://towardsdatascience.com/the-subtle-art-of-fixing-and-modifying-learning-rate-f1e22b537303>

Heuristic for learning rate




- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem
- learning rate decay



Heuristic for learning rate



- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem
- learning rate decay



Normalization can help



- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem
- **normalization is always helpful** when there are hyper-parameters

$$\begin{array}{llll}
 f(x) = x^2 & \longrightarrow & \nabla f(x) = 2x & \longrightarrow & x^{k+1} = x^k + 2t^k x^k & t^* = 1 \\
 \downarrow y = 10x & & & & \searrow y = 10x & \\
 g(y) = \frac{y^2}{100} & \longrightarrow & \nabla g(y) = \frac{y}{50} & \longrightarrow & y^{k+1} = y^k + 2t^k y^k & \\
 & & & & & s^* = 100
 \end{array}$$



Batch normalization



- (S)GD needs line search
 - one may converge to a region where the gradient's magnitude is small, and then stay there for a very long time
 - adjustable step-length can avoid the situation, but not feasible for large-scale problem
- **normalization is always helpful** when there are hyper-parameters
 - Batch Normalizing: normalize the batch of the input data of neural networks to zero-mean and constant standard deviation



RMSProp



- Root Mean Square Propagation

divide the learning rate for a weight by a running average of the magnitudes of recent gradients for that weight

$$z_{k+1} = \rho z_k + (1 - \rho)(\nabla f(x_k) \cdot \nabla f(x_k))$$

$$x_{k+1} = x_k - \frac{\eta}{\sqrt{z_{k+1} + \varepsilon}} \nabla f(x_k)$$

also for stochastic version

- adaptive learning rate

AdaGrad



- Adaptive gradient algorithm

improves convergence performance, when data are sparse and sparse parameters are more informative

$$G_{ii} = \sum_{t=1}^k (\nabla f(x_t)(i))^2$$



$$x_{k+1} = x_k - \frac{\eta}{\sqrt{G_{ii} + \varepsilon}} \nabla f(x_k)$$

- overall time (AdaGrad)/sliding window (RMSProp)

Adam (Adaptive Moment Estimation)



- combines the heuristics of

- momentum

- RMSProp

$$m_{k+1} = \rho m_k + (1 - \rho) \nabla f(x_k) \quad \text{💬}$$

$$v_{k+1} = \rho v_k + (1 - \rho) (\nabla f(x_k) \cdot \nabla f(x_k)) \quad \text{💬}$$


$$\hat{m}_{k+1} = m_{k+1} / (1 - \beta_1^{k+1})$$

$$\hat{v}_{k+1} = v_{k+1} / (1 - \beta_2^{k+1})$$

$$x_{k+1} = x_k - \eta \hat{m}_{k+1} / (\sqrt{\hat{v}_{k+1}} + \varepsilon)$$




Example






Example



Example



Shuffling and Curriculum Learning



- how to choose the batch for SGD

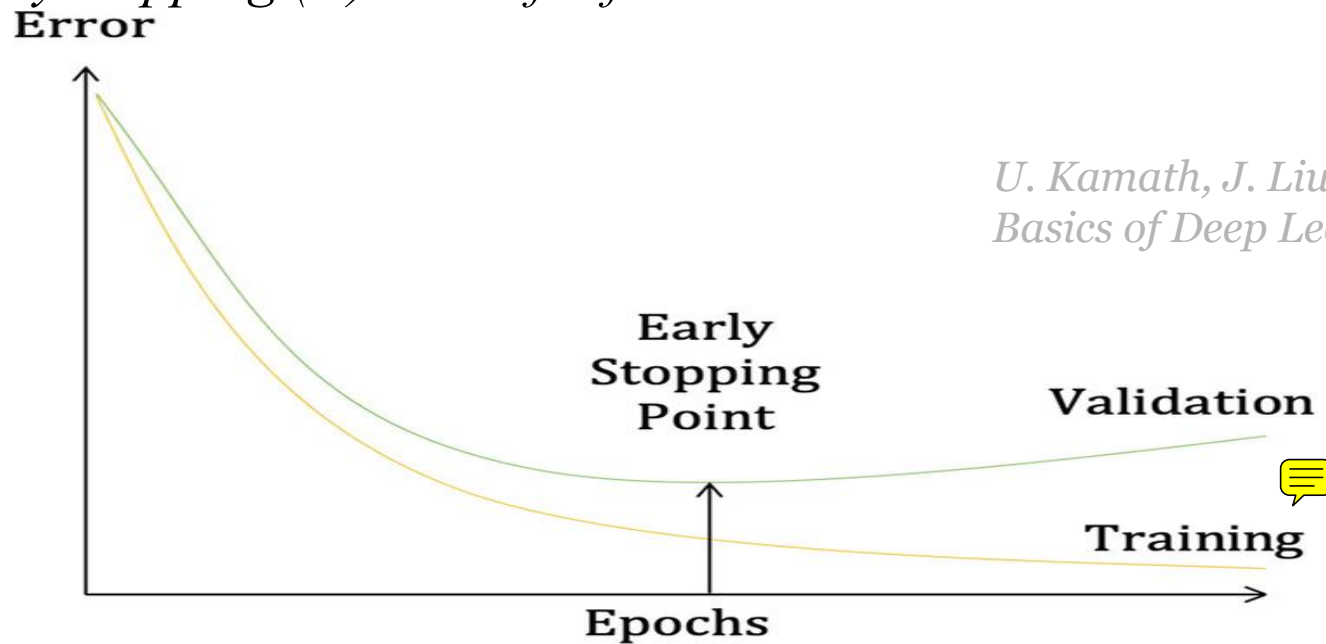
$$x^{k+1} = x^k - t \sum_{i \in I} \nabla f_i(x)$$

- shuffling: pure randomness to avoid biasness
- curriculum: supplying the training examples in a meaningful order may actually lead to improved performance

Early Stop



- “*Early stopping (is) beautiful free lunch*”




*U. Kamath, J. Liu, J. Whitaker:
Basics of Deep Learning*

- could be regarded as a regularization term
- especially, a few-shot learning by fine-tune

Early Stop



- “Early stopping (is) beautiful free lunch”



*U. Kamath, J. Liu, J. Whitaker:
Basics of Deep Learning*

- could be regarded as a regularization term
- especially, a few-shot learning by fine-tune

Non-smoothness



$$\min_x x^\top x + C \sum_{i=1}^m \max\{0, 1 - b_i(x^\top a_i)\}$$

- Pegasos: Primal Estimated sub-GrAdient SOLver for SVM

INPUT: S, λ, T, k

INITIALIZE: Choose \mathbf{w}_1 s.t. $\|\mathbf{w}_1\| \leq 1/\sqrt{\lambda}$

FOR $t = 1, 2, \dots, T$

Choose $A_t \subseteq S$, where $|A_t| = k$

Set $A_t^+ = \{(\mathbf{x}, y) \in A_t : y \langle \mathbf{w}_t, \mathbf{x} \rangle < 1\}$

Set $\eta_t = \frac{1}{\lambda t}$

Set $\mathbf{w}_{t+\frac{1}{2}} = (1 - \eta_t \lambda) \mathbf{w}_t + \frac{\eta_t}{k} \sum_{(\mathbf{x}, y) \in A_t^+} y \mathbf{x}$

Set $\mathbf{w}_{t+1} = \min \left\{ 1, \frac{1/\sqrt{\lambda}}{\|\mathbf{w}_{t+\frac{1}{2}}\|} \right\} \mathbf{w}_{t+\frac{1}{2}}$

OUTPUT: \mathbf{w}_{T+1}

Theorem 3. Assume that the conditions stated in Thm. 2 hold. Let $\delta \in (0, 1)$. Then, with probability of at least $1 - \delta$ over the choices of (A_1, \dots, A_T) and the index r we have that

$$f(\mathbf{w}_r) \leq f(\mathbf{w}^*) + \frac{c \ln(T)}{\delta \lambda T}.$$

Stochastic in the dual



- Dual coordinate descent
- Stochastic dual coordinate descent
- An example

Stochastic Dual Coordinate Ascent Methods for Regularized Loss Minimization

Shai Shalev-Shwartz

*Benin school of Computer Science and Engineering
The Hebrew University
Jerusalem, 91904, Israel*

SHAIS@CS.HUJI.AC.IL

Tong Zhang

*Department of Statistics
Rutgers University
Piscataway, NJ, 08854, USA*

TZHANG@STAT.RUTGERS.EDU

$$\begin{aligned} \min_{\mathbf{x}} \quad & P(\mathbf{x}) \triangleq \mu \|\mathbf{x}\|_1 + \frac{1}{m} \sum_{i=1}^m L_{\tau,c}(-y_i(\mathbf{u}_i^\top \mathbf{x})) \\ \text{s.t.} \quad & \|\mathbf{x}\|_2 \leq 1. \end{aligned}$$



$$\begin{aligned} \text{maximize}_{\mathbf{s}, \mathbf{t}} \quad & D(\mathbf{s}, \mathbf{t}) \triangleq c \sum_{i=1}^m t_i - \left\| \sum_{i=1}^m t_i y_i \mathbf{u}_i - \mathbf{s} \right\|_2 \\ \text{s.t.} \quad & \|\mathbf{s}\|_\infty \leq \mu, \quad -\frac{\tau}{m} \leq \mathbf{t} \leq \frac{1}{m}. \end{aligned}$$

- primal space: non-smooth, constraint
- dual space: smooth, separable constraint
- randomly select coordinate and update

Non-convex



- Nowadays, you may meet many non-convex problems
 - finding a global optimum is NP hard
 - do not try theoretical analysis, although it is an open problem
 - randomness may help but not efficient
 - global optimality cannot be guaranteed,
by neither GD nor SGD
 - some arguments: stochastic helps global optimization

Non-convex



- global optimum \rightarrow local optimum \rightarrow saddle points \rightarrow convergence
 - for non-convex functions, we can still have
 - Second-differentiable condition
 - Lipschitz-continuous condition
 - noise has bounded variance

- but we do not have convexity

$$f^* \geq f(x) - \frac{1}{2m} \|\nabla f(x)\|_2^2$$

- non-convex SGD converges, but does not necessarily go to a saddle point, let alone a local optimum

1

Stochastic Gradient Descent

2

SGD with Momentum

3

Limited-memory BFGS

4

Coordinate Update



Estimate the second-order information



- estimate the second-order information by several gradient
 - Hessian: the change (gradient) of gradient
 - $\nabla f(x^k) \approx \nabla f(x^{k-1}) + t \nabla^2 f(x^{k-1})$
- using momentum information to lookahead gradient

$$z_{k+1} = x_k - t_k \nabla f(x_k)$$

$$x_{k+1} = x_k + \delta_k (x_k - x_{k-1}) \quad \delta_k \in [0,1)$$



- using a matrix to approach the Hessian matrix, *Quasi-Newton method*



$$B_k \Delta x_k = -\nabla f(x_k)$$

$$\Delta x_k = -B_k^{-1} \nabla f(x_k)$$

key problem: how to update B or B^{-1} to approach Hessian or its inverse

Construct the inverse



$$q_k = \nabla f(x_{k+1}) - \nabla f(x_k) \approx \nabla^2 f(x_k)(x_{k+1} - x_k)$$

for quadratic case



$$q_k = \nabla f(x_{k+1}) - \nabla f(x_k) = \nabla^2 f(x_k)(x_{k+1} - x_k)$$

- let $H = (\nabla^2 f(x_k))^{-1}$, there should be $Hq_k = x_{k+1} - x_k$
- we want to use rank one correction update

$$H_{k+1} = H_k + a_k z_k z_k^\top$$

$$H_{k+1}q_k = x_{k+1} - x_k$$



$$H_{k+1} = H_k + \frac{(x_{k+1} - x_k - H_k q_k)(x_{k+1} - x_k - H_k q_k)^\top}{q_k^\top (x_{k+1} - x_k - H_k q_k)}$$

$$\Delta x^k = -H_{k+1} \nabla f(x_k) \quad \text{modified Newton's method with rank 1 correction}$$

DFP and BFGS method



- Davidon-Fletcher-Powell uses Rank two correction

$$H_{k+1} = H_k + \underbrace{a_k z_k z_k^\top + \beta_k y_k y_k^\top}_{\text{Rank two correction}} \quad H_{k+1} q_k = x_{k+1} - x_k$$

$$H_{k+1} = H_k + \frac{(x_{k+1} - x_k)(x_{k+1} - x_k)^\top}{(x_{k+1} - x_k)^\top q_k} - \frac{H_k q_k q_k^\top H_k}{q_k^\top H_k q_k}$$

- Broyden-Fletcher-Goldfarb-Shanno estimate the Hessian instead of its inverse

$$H_{k+1} q_k = x_{k+1} - x_k \quad \nabla^2 f(x_k)(x_{k+1} - x_k) = q_k$$

by symmetry, you could directly change the position

$$Q_{k+1} = Q_k + \frac{(q_k - H_k(x_{k+1} - x_k))(q_k - H_k(x_{k+1} - x_k))^\top}{(x_{k+1} - x_k)^\top (p_k - H_k(x_{k+1} - x_k))}$$

DFP and BFGS method



- Davidon-Fletcher-Powell uses Rank two correction

$$H_{k+1} = H_k + \underbrace{a_k z_k z_k^\top + \beta_k y_k y_k^\top}_{H_{k+1} q_k = x_{k+1} - x_k}$$

$$H_{k+1} = H_k + \frac{(x_{k+1} - x_k)(x_{k+1} - x_k)^\top}{(x_{k+1} - x_k)^\top q_k} - \frac{H_k q_k q_k^\top H_k}{q_k^\top H_k q_k}$$

- Broyden-Fletcher-Goldfarb-Shanno estimate the Hessian instead of its inverse

$$H_{k+1} q_k = x_{k+1} - x_k \quad \nabla^2 f(x_k)(x_{k+1} - x_k) = q_k$$

by symmetry, you could directly change the position

$$Q_{k+1} = Q_k + \frac{q_k q_k^\top}{(x_{k+1} - x_k)^\top q_k} - \frac{Q_k (x_{k+1} - x_k)(x_{k+1} - x_k)^\top Q_k}{(x_{k+1} - x_k)^\top Q_k (x_{k+1} - x_k)}$$

BFGS algorithm



- for optimization, we actually need Q^{-1}

Sherman-Morrison matrix inversion formula

$$(A + uv^T)^{-1} = A^{-1} - \frac{A^{-1}uv^T A^{-1}}{1 + v^T A^{-1}u}$$

- let $B_{k+1} = Q_{k+1}^{-1}$ and apply Sherman-Morrison formula twice

$$B_{k+1} = B_k - \frac{B_k(x_{k+1} - x_k)^T B_k}{(x_{k+1} - x_k)^T B_k (x_{k+1} - x_k)} + \frac{q_k q_k^T}{q_k^T (x_{k+1} - x_k)}$$

BFGS is almost the same as DFP, but in practice, BFGS seems work better
—— Tibshirani

Limited-memory BFGS



- what we really need is $B_{k+1} \nabla f(x_{k+1})$
- instead of storing B in BFGS

$$B_{k+1} = B_k - \frac{B_k(x_{k+1} - x_k)^\top B_k}{(x_{k+1} - x_k)^\top B_k(x_{k+1} - x_k)} + \frac{q_k q_k^\top}{q_k^\top (x_{k+1} - x_k)}$$



- we only need to calculate $B_{k+1} \nabla f(x_{k+1})$
 - maintain two recursive sequences of the past m updates for x and $\nabla f(x)$
 - starting from $B_{k-m} = I$
- modification/improvement
 - L-BFGS-b: to handle box constraints
 - O-LBFGS: online version with randomly drawn subset
 - forgetting/mini-batch/....

1

Stochastic Gradient Descent

2

SGD with Momentum

3

Limited-memory BFGS

4

Coordinate Update



Parallel computing I




$$\min_x \|B - Ax\|_2^2$$


- gradient descent

$$\begin{aligned} x^{k+1} &= x^k + \lambda_k A^\top (B - Ax^k) \\ &= (I - \lambda_k A^\top A) x^k + \lambda_k A^\top B \\ &\triangleq \underbrace{T_k}_{\text{red circle}} x^k + b_k \end{aligned}$$

- coordinate update
 - synchronous strategy
 - asynchronous strategy



Synchronous
(wait for the slowest)



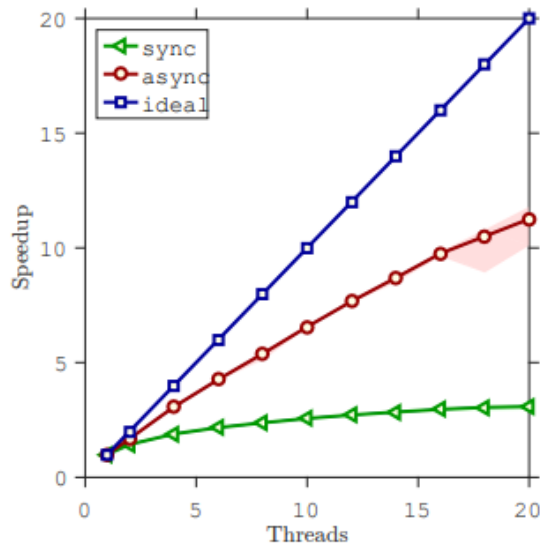
Asynchronous
(non-stop, no wait)

Z. Peng, Y. Xu, M. Yan, W. Yin, ARock: an algorithmic framework for asynchronous parallel coordinate updates, *SIAM Journal on Scientific Computing*, 2016

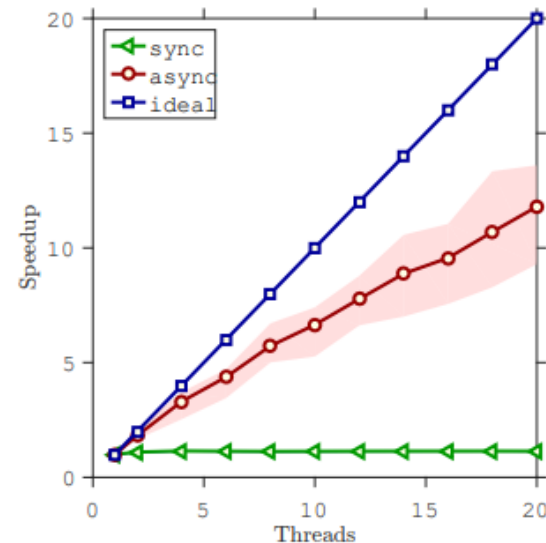
Parallel computing I



- convergence for synchronous strategy has no problem
- asynchronous strategy is much more efficient



dataset "news20"



dataset "url"

Parallel computing I



- convergence for synchronous strategy has no problem
- asynchronous strategy is much more efficient
- convergence is a problem

Z. Peng, Y. Xu, M. Yan, W. Yin,
ARock: an algorithmic framework
for asynchronous parallel
coordinate updates, *SIAM Journal
on Scientific Computing*, 2016

notation:

- $m = \#$ coordinates
- $\tau =$ the maximum delay
- uniform selection $p_i \equiv \frac{1}{m}$ (just for simplicity)

Theorem (almost sure convergence)

Assume that T is nonexpansive and has a fixed point. Use step sizes $\eta_k \in [\epsilon, \frac{1}{2m-1/2\tau+1})$, $\forall k$. Then, with probability one, $x^k \rightarrow x^* \in \text{Fix}T$.

Parallel computing I



- proof

- **typical inequality:**

$$\begin{aligned}\|x^{k+1} - x^*\|^2 &\leq \|x^k - x^*\|^2 - c\|Tx^k - x^k\|^2 \\ &\quad + \text{harmful terms}(x^{k-1}, \dots, x^{k-\tau})\end{aligned}$$

- **descent inequality under a new metric:**

$$\mathbb{E} \left(\|\mathbf{x}^{k+1} - \mathbf{x}^*\|_M^2 \mid \mathcal{X}^k \right) \leq \|\mathbf{x}^k - \mathbf{x}^*\|_M^2 - c \|Tx^k - x^k\|^2$$

where

- $\mathbf{x}^k = (x^k, x^{k-1}, \dots, x^{k-\tau}) \in \mathcal{H}^{\tau+1}, k \geq 0$
- $\mathbf{x}^* = (x^*, x^*, \dots, x^*) \in \mathbf{X}^* \subseteq \mathcal{H}^{\tau+1}$
- M is a positive definite matrix.
- c depends on (η_k, m, τ)

Z. Peng, Y. Xu, M. Yan, W. Yin,
ARock: an algorithmic framework
for asynchronous parallel
coordinate updates, *SIAM Journal
on Scientific Computing*, 2016

Parallel computing I



- proof

- **typical inequality:**

$$\begin{aligned}\|x^{k+1} - x^*\|^2 &\leq \|x^k - x^*\|^2 - c\|Tx^k - x^k\|^2 \\ &\quad + \text{harmful terms}(x^{k-1}, \dots, x^{k-\tau})\end{aligned}$$

- **descent inequality under a new metric:**

$$\mathbb{E} \left(\|\mathbf{x}^{k+1} - \mathbf{x}^*\|_M^2 \mid \mathcal{X}^k \right) \leq \|\mathbf{x}^k - \mathbf{x}^*\|_M^2 - c \|Tx^k - x^k\|^2$$

where

- $\mathbf{x}^k = (x^k, x^{k-1}, \dots, x^{k-\tau}) \in \mathcal{H}^{\tau+1}, k \geq 0$
- $\mathbf{x}^* = (x^*, x^*, \dots, x^*) \in \mathbf{X}^* \subseteq \mathcal{H}^{\tau+1}$
- M is a positive definite matrix.
- c depends on (η_k, m, τ)

Z. Peng, Y. Xu, M. Yan, W. Yin,
ARock: an algorithmic framework
for asynchronous parallel
coordinate updates, *SIAM Journal
on Scientific Computing*, 2016

THANKS

