

Ve492: Introduction to Artificial Intelligence

Hidden Markov Models II



Paul Weng


UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

Today

- ❖ Quick Review of (H)MM
- ❖ Particle Filtering
- ❖ Viterbi Algorithm for Most Likely Explanation
- ❖ Dynamic Bayesian Network


Markov Chain




- ❖ You know transition probabilities, $P(X_t | X_{t-1})$, and $P(X_4)$.
- ❖ Write an equation to compute $P(X_5)$.

$$\begin{aligned} P(X_5) &= \sum_{x_4} P(x_4, X_5) \\ &= \sum_{x_4} P(X_5 | x_4) P(x_4) \end{aligned}$$

- ❖ More generally: $P_{t+1} = T^T P_t$ where $P_t = P(X_t)$
- ❖ Stationary distribution: $P_\infty = T^T P_\infty$



HMM, Filtering, Forward Algorithm



Filtering: What is the current state, given all evidence?

$$P(X_{t+1} | e_{1:t+1}) = \alpha \underbrace{P(e_{t+1} | X_{t+1})}_{\text{Update}} \underbrace{\sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t})}_{\text{Predict}}$$

Normalize

Update


Predict

$$f_{1:t+1} = \text{FORWARD}(f_{1:t}, e_{t+1})$$

Forward Algorithm

Query: What is the current state, given all of the current and past evidence?


Marching **forward** through the HMM network



Forward Algorithm

Query: What is the current state, given all of the current and past evidence?


Marching **forward** through the HMM network



Forward Algorithm

Query: What is the current state, given all of the current and past evidence?


Marching **forward** through the HMM network




Forward Algorithm

Query: What is the current state, given all of the current and past evidence?

Marching **forward** through the HMM network



Particle Filtering



Quiz: Algorithms for Filtering

- ❖ Which algorithms can be applied to the filtering task?
 - ❖ Variable elimination
 - ❖ Prior sampling
 - ❖ Rejection sampling
 - ❖ Likelihood weighting
 - ❖ Gibbs sampling

Limitations of Current Algorithms

- ❖ Exact inference infeasible if state space is large or continuous
 - ❖ e.g., when $|X|$ is more than 10^6 or so (tracking 3 ghosts in a 10x20 world)
 - ❖ e.g., robot localization
- ❖ Previous sampling methods do not exploit the temporal structure of an HMM

Particle Filtering

❖ PF = approximate inference for filtering task

❖ Principle of Particle Filtering

- ❖ Track samples of X_t , not all values
- ❖ Samples are called particles
- ❖ Time per step is linear in the number of samples
- ❖ But: number needed may be large
- ❖ In memory: list of particles, not states


0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5



	●	
		● ●
	● ●	● ● ● ●

Representation: Particles

- ❖ Our representation of $P(X)$ is now a list of N particles (samples)
 - ❖ Generally, $N \ll |X|$
 - ❖ Storing map from X to counts would defeat the point
- ❖ $P(x)$ approximated by number of particles with value x
 - ❖ So, many x may have $P(x) = 0$!
 - ❖ More particles, more accuracy
- ❖ For now, all particles have a weight of 1
- ❖ In PF, a set of samples approximates $f_{1:t}(X_t) = P(X_t | e_{1:t})$



Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particle Filtering: Elapse Time


- ❖ This first step captures the passage of time
- ❖ Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- ❖ This is like prior sampling – samples' frequencies reflect the transition probabilities
 - ❖ Here, most samples move clockwise, but some move in another direction or stay in place
- ❖ New particles approximates: $\sum_{x_t} P(X_{t+1} | x_t) f_{1:t}(x_t)$
- ❖ **Consistency:** If enough samples, close to exact values before and after


Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particle Filtering: Observe

❖ Slightly trickier:

- ❖ Don't sample observation, fix it
- ❖ Similar to likelihood weighting, weight samples based on the evidence


$$w(x) = P(e | x)$$

$$f_{1:t+1}(X) \propto w(X) \sum_{x_t} P(X | x_t) f_{1:t}(x_t)$$

- ❖ As before, the probabilities don't sum to one, since all have been weighted
- ❖ What do they sum to?


Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4




Particle Filtering: Resample

- ❖ Rather than tracking weighted samples, we resample
- ❖ N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- ❖ This is equivalent to renormalizing the distribution
- ❖ Now the update is complete for this time step, continue with the next one


Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$




(New)
Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)

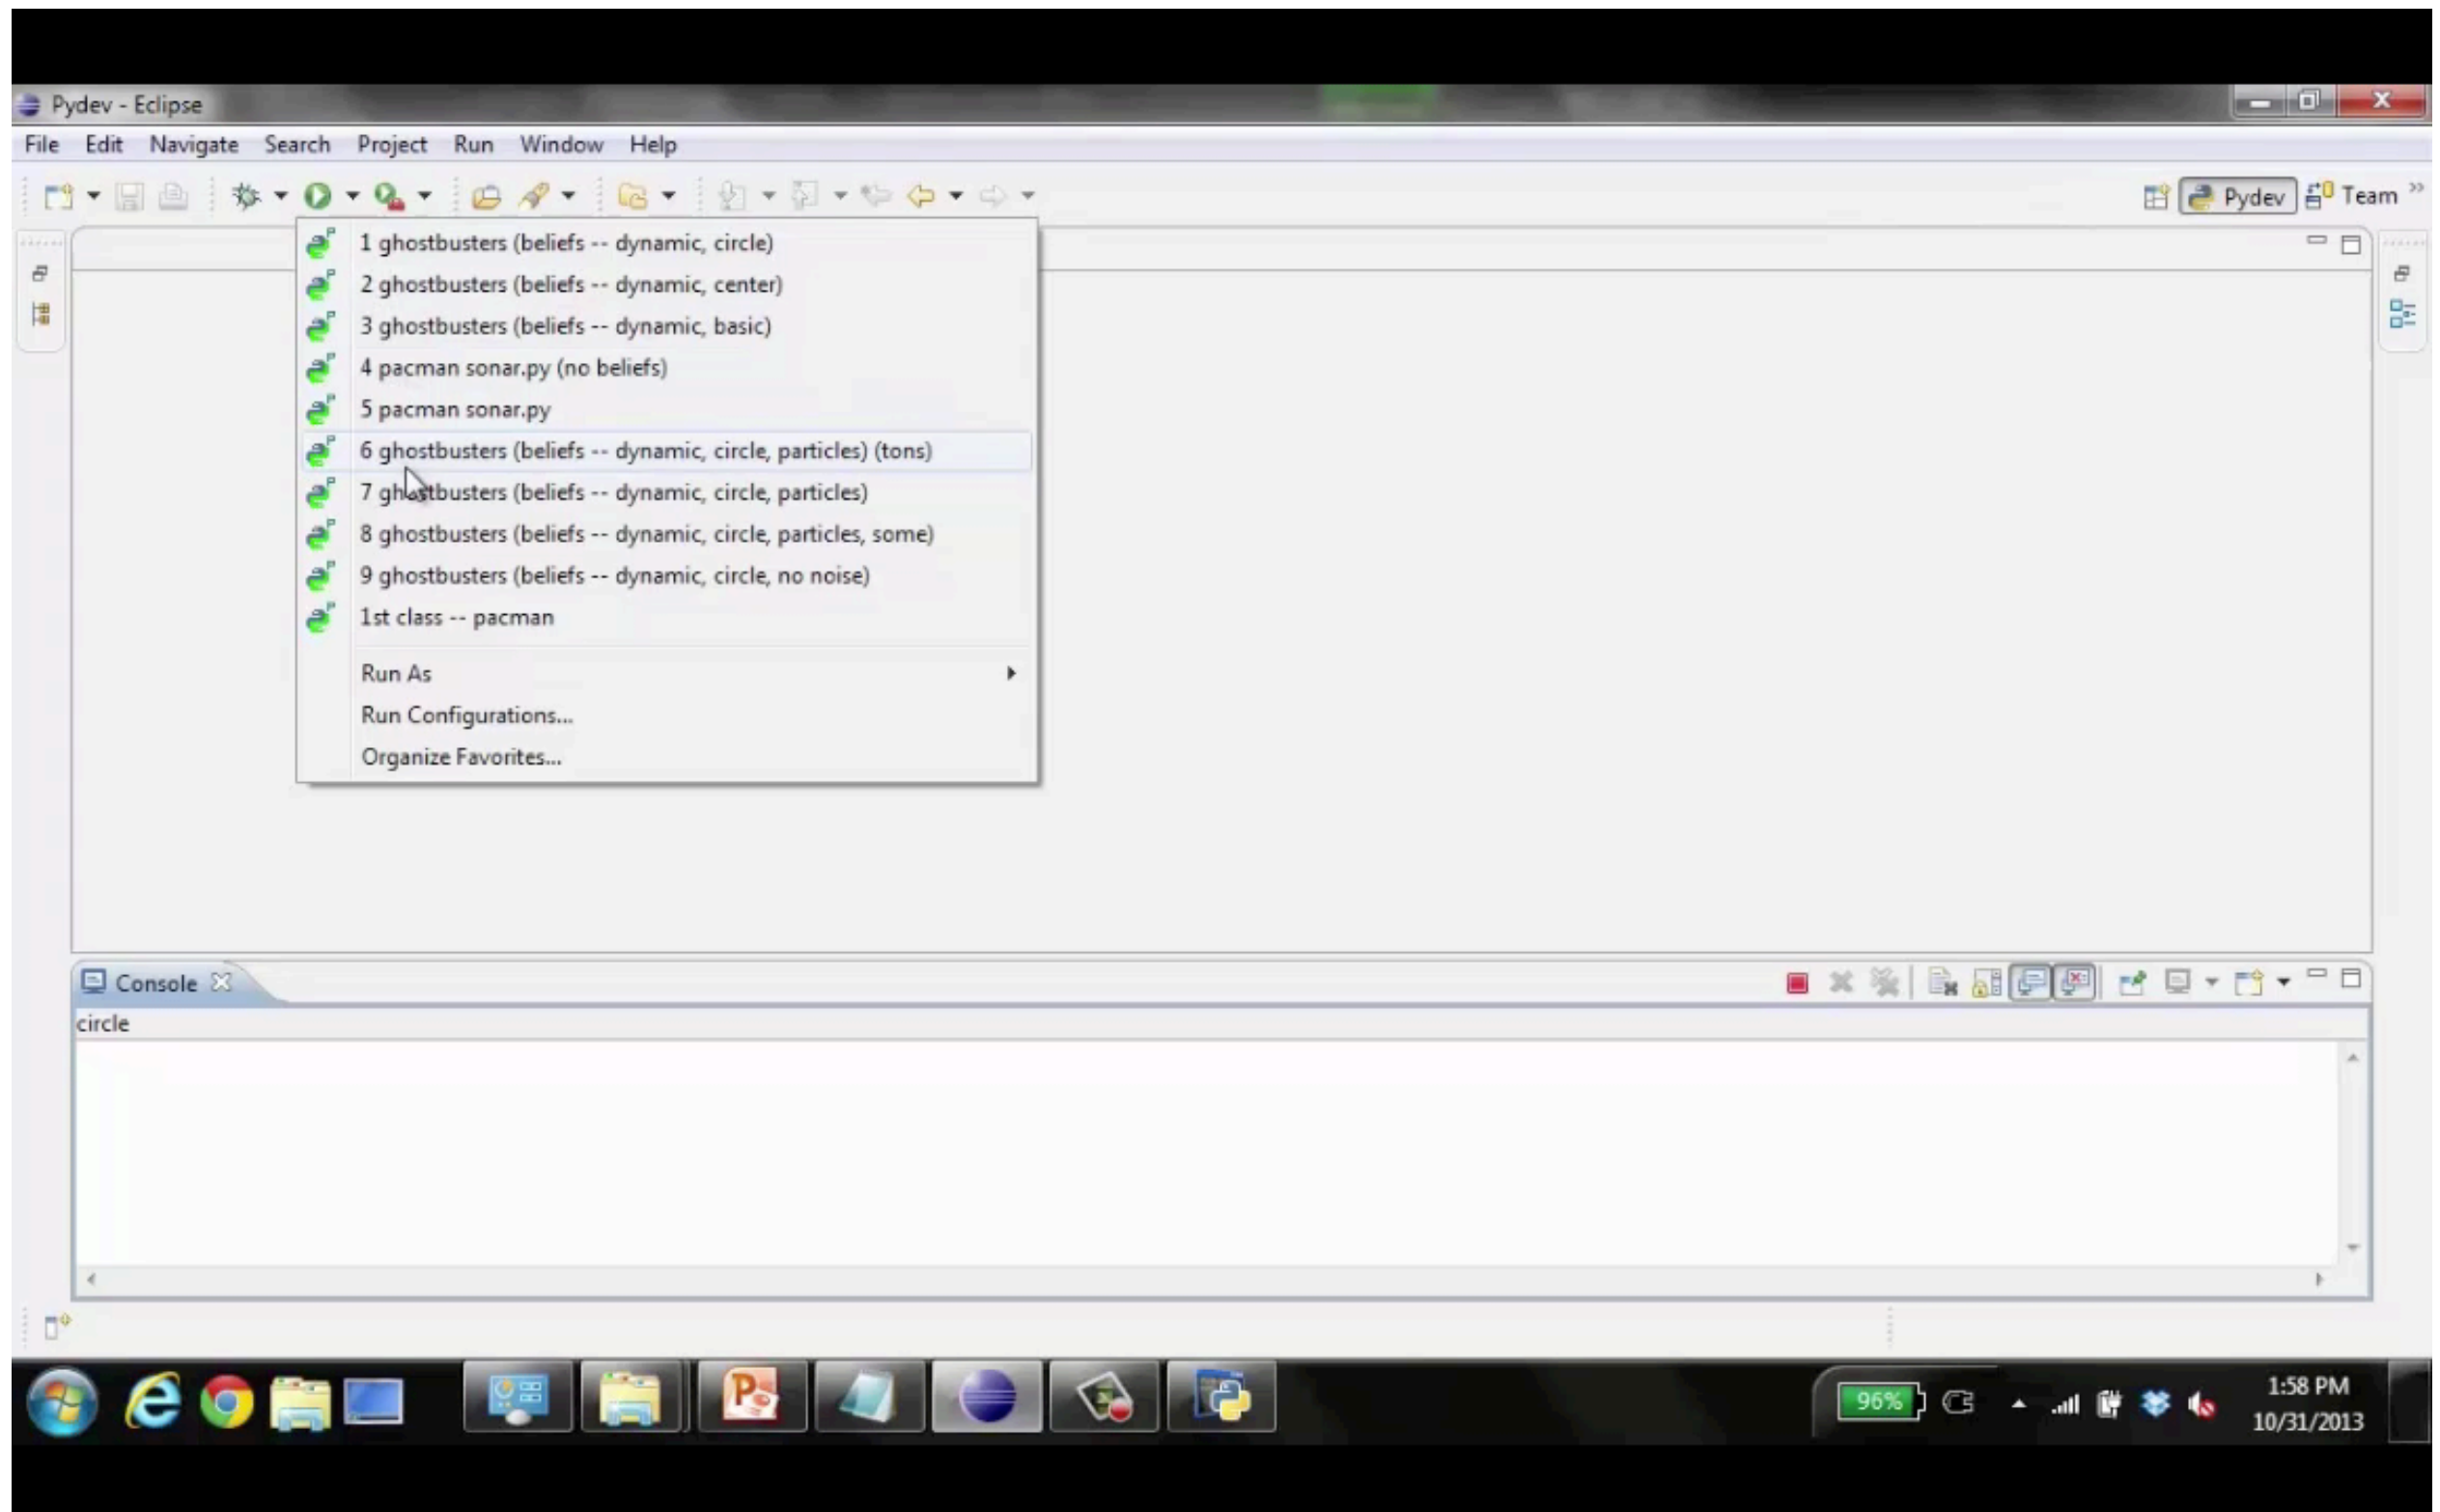


Summary: Particle Filtering

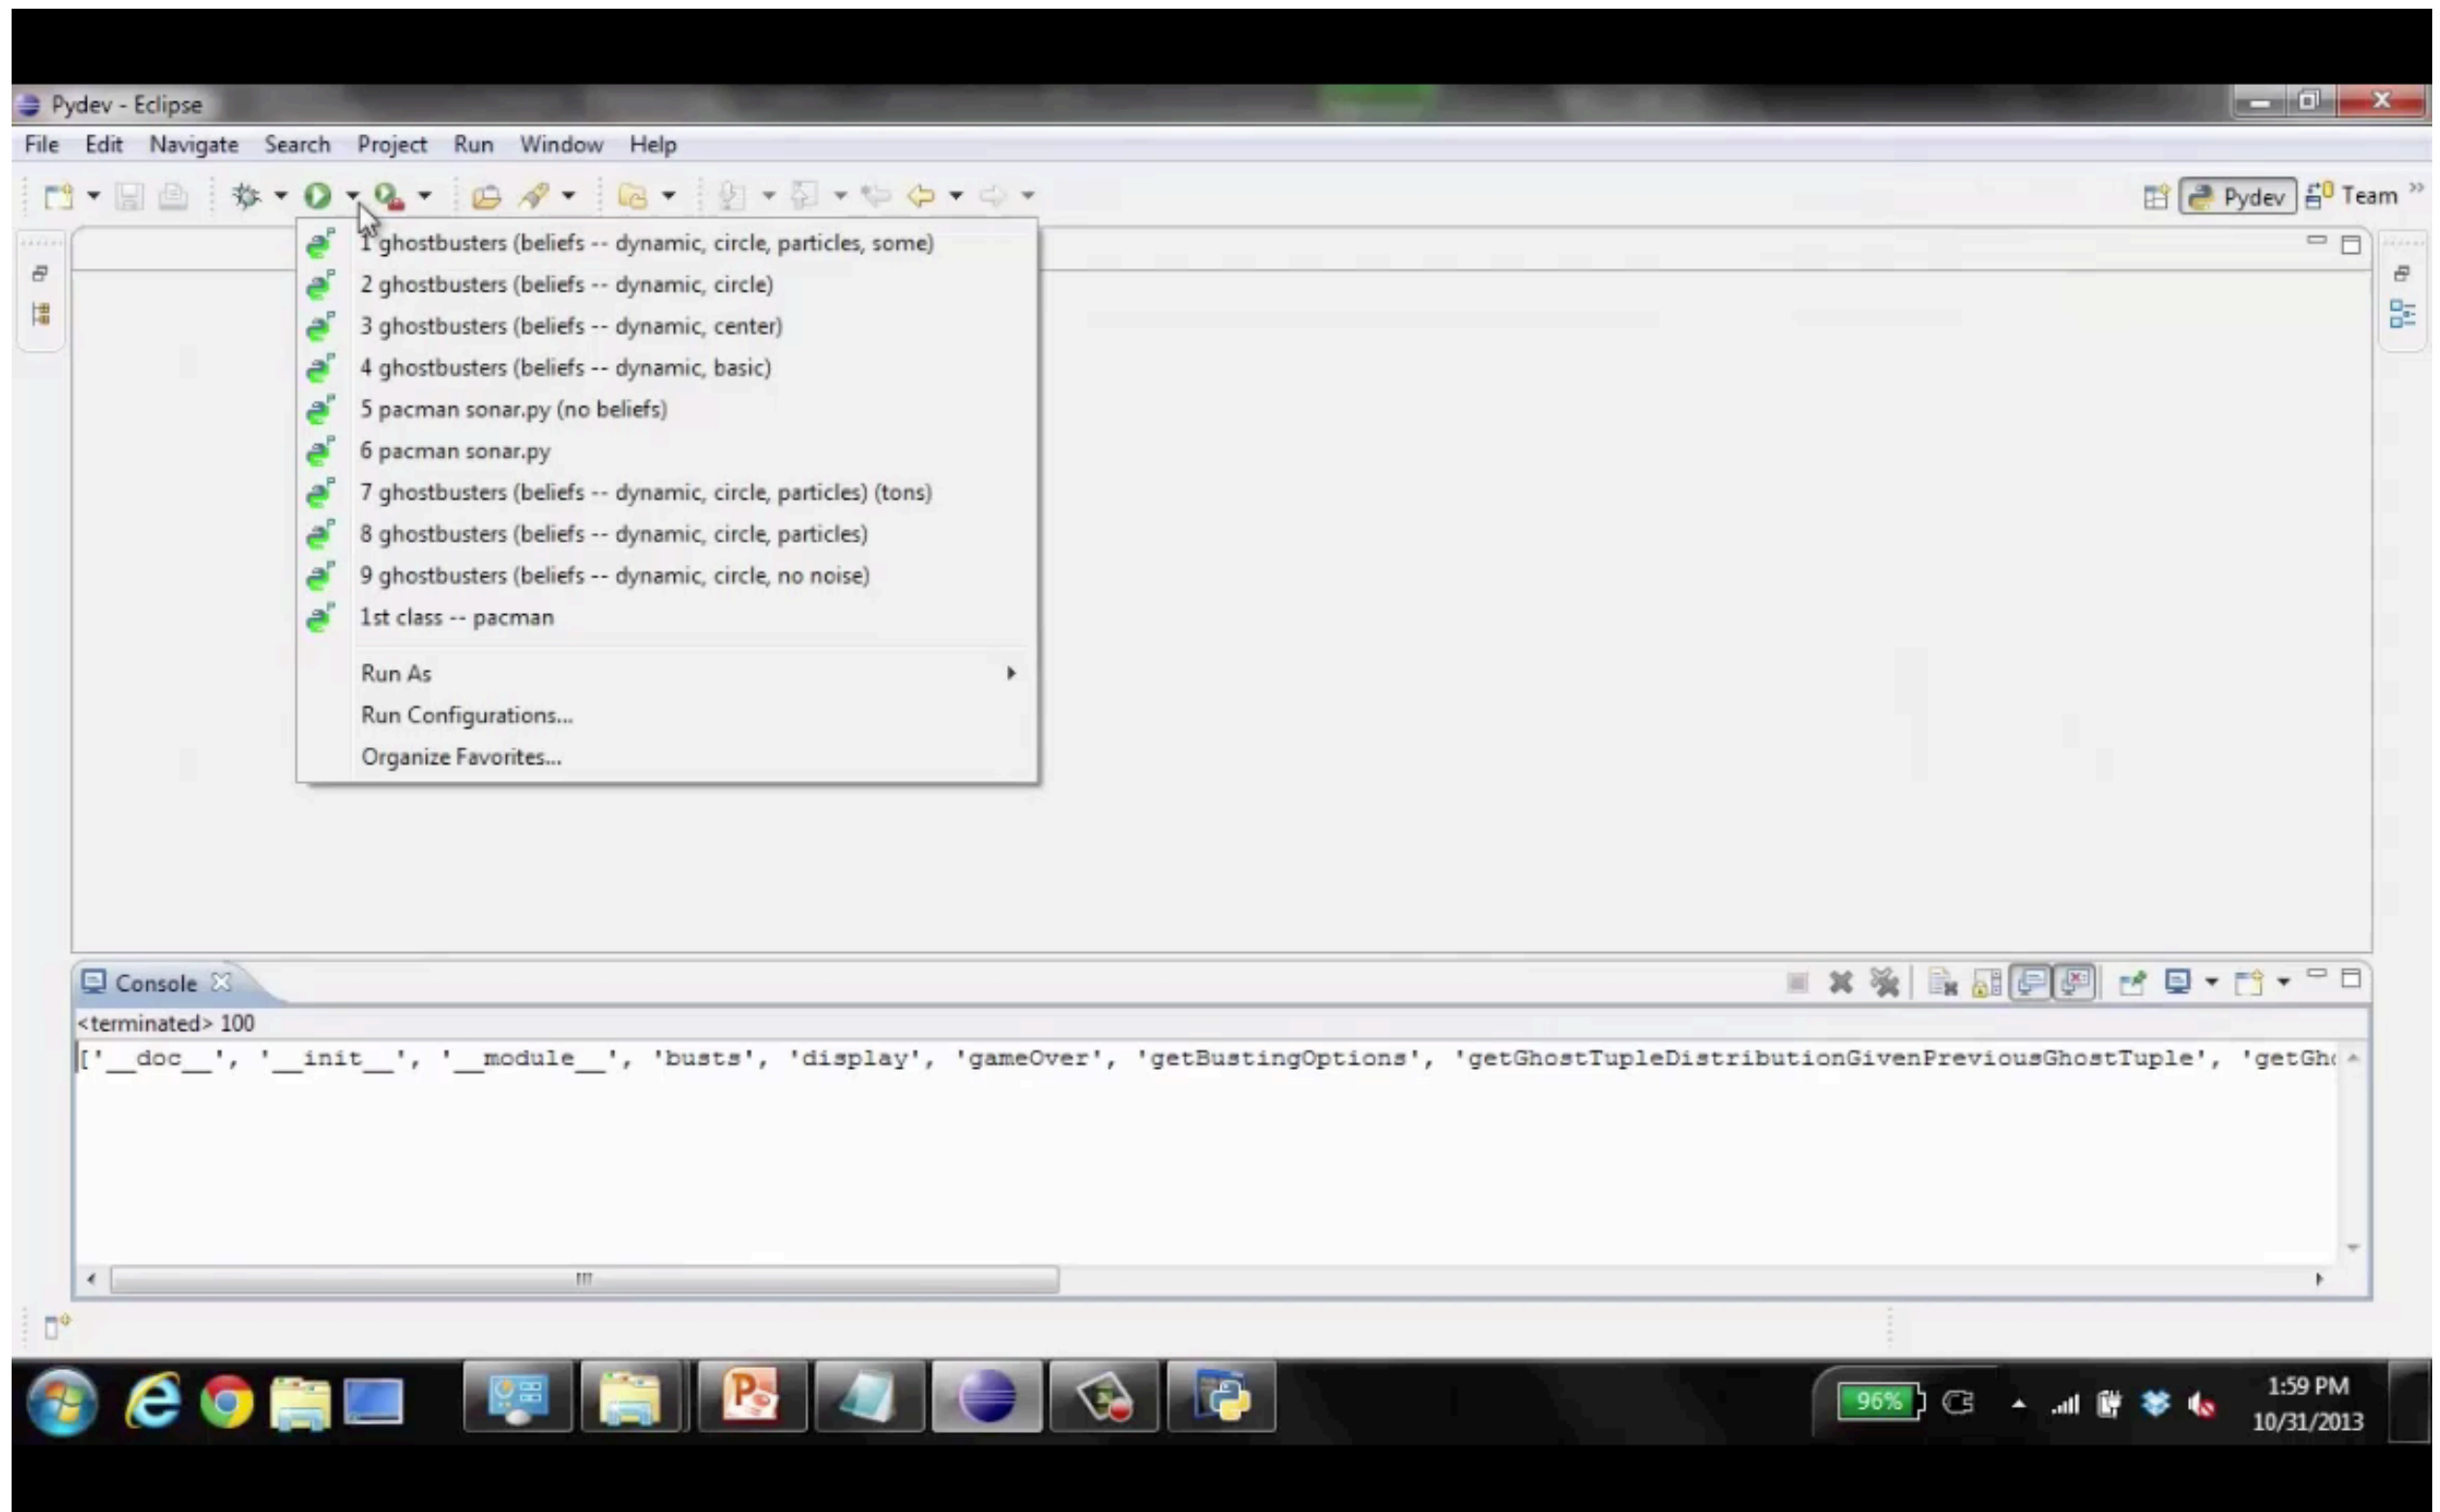
- ❖ **Particles:** track samples of states rather than an explicit distribution



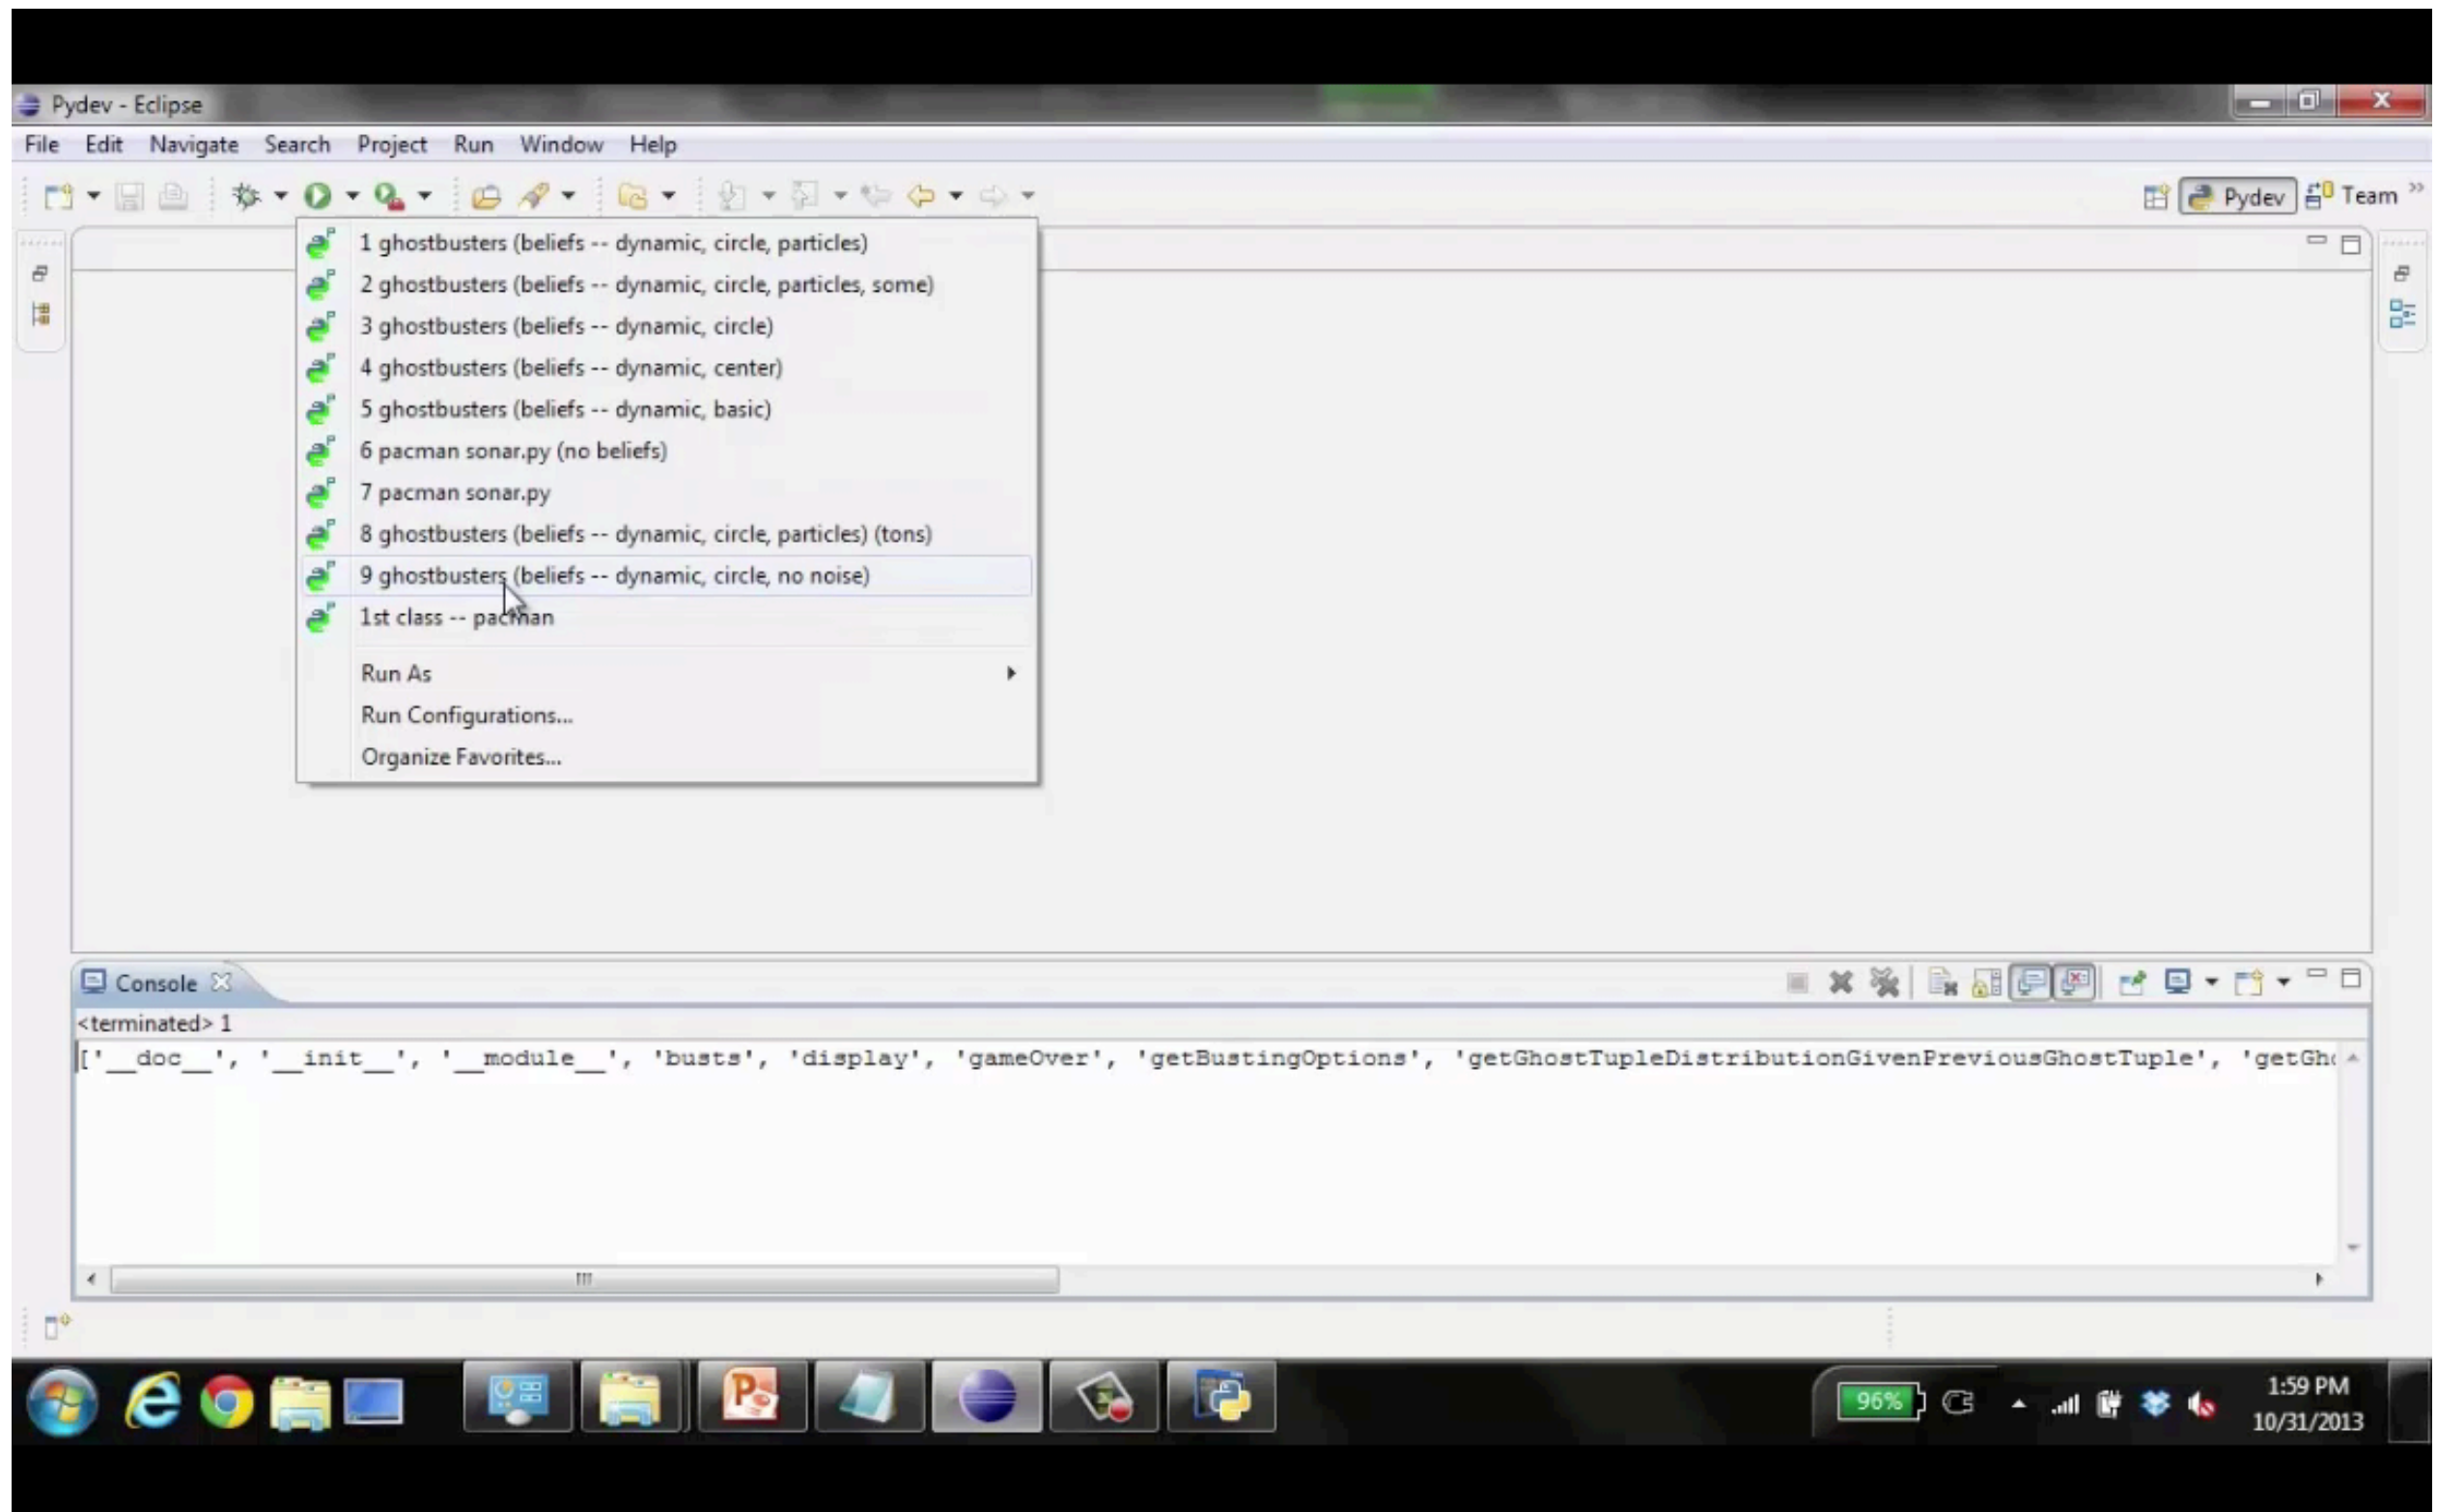
Video of Demo – Moderate Number of Particles



Video of Demo – One Particle





Video of Demo – Huge Number of Particles




Robot Localization

In robot localization:

- ❖ We know the map, but not the robot's position
- ❖ Observations may be vectors of range finder readings
- ❖ State space and readings are typically continuous
- ❖ Particle filtering is a main technique



Particle Filter Localization (Laser)



[Dieter Fox, et al.]

Most Likely Explanation




Inference tasks

- ❖ **Filtering**: $P(X_t | e_{1:t})$
 - ❖ **belief state**—input to the decision process of a rational agent
- ❖ **Prediction**: $P(X_{t+k} | e_{1:t})$ for $k > 0$
 - ❖ evaluation of possible action sequences; like filtering without the evidence
- ❖ **Smoothing**: $P(X_k | e_{1:t})$ for $0 \leq k < t$
 - ❖ better estimate of past states, essential for learning
- ❖ **Most likely explanation**: $\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$
 - ❖ speech recognition, decoding with a noisy channel

HMMs: MLE Queries


- ❖ HMMs defined by
 - ❖ States X
 - ❖ Observations E
 - ❖ Initial distribution: $P(X_1)$
 - ❖ Transitions: $P(X_t | X_{t-1})$
 - ❖ Emissions: $P(E | X)$



- ❖ New query: most likely explanation: $\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$
- ❖ New method: the Viterbi algorithm


State Trellis

- ❖ State trellis: graph of states and transitions over time



- ❖ Each arc represents some transition $x_{t-1} \rightarrow x_t$
- ❖ Each arc has weight $P(x_t|x_{t-1})P(e_t|x_t)$
- ❖ Each path is a sequence of states
- ❖ Product of weights on a path = sequence's probability along with the evidence
- ❖ Forward algorithm computes sums of paths, Viterbi computes best paths

Forward / Viterbi algorithms



Forward Algorithm (sum)

For each state at time t , keep track of the **total probability of all paths** to it


$$\begin{aligned} \mathbf{f}_{1:t+1} &= \text{FORWARD}(\mathbf{f}_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1}|X_{t+1}) \sum_{x_t} P(X_{t+1}|x_t) \mathbf{f}_{1:t} \end{aligned}$$

Viterbi Algorithm (max)

For each state at time t , keep track of the **maximum probability of any path** to it

$$\begin{aligned} \mathbf{m}_{1:t+1} &= \text{VITERBI}(\mathbf{m}_{1:t}, e_{t+1}) \\ &= P(e_{t+1}|X_{t+1}) \max_{x_t} P(X_{t+1}|x_t) \mathbf{m}_{1:t} \end{aligned}$$

Why is This True?



$$m_1[x_1] = P(e_1|x_1)P(x_1)$$

probability of best path $1 : t$ that ends at x_t


$$m_2[x_2] = \max\{P(e_2|x_2)P(x_2|x_1 = \text{rain})m_1[x_1 = \text{rain}], P(e_2|x_2)P(x_2|x_1 = \text{sun})m_1[x_1 = \text{sun}]\}$$

$$= \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1]$$

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

Now What?



$m_N[x_N]$ probability of best path $1 : N$ that ends at x_N


what is the last state on the most likely path?

$$\arg \max_{x_N} m_N[x_N]$$

what is the *second to last* state on the most likely path?

A Tricky Counter-Example

$$P(e_1|x_1)P(x_1) \quad P(e_2|x_2)P(x_2|x_1)$$



$$\arg \max_{x_1} m_1[x_1] = ?$$

sun!

$$m_2[x_2] = \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1]$$

$$m_2[x_2 = \text{sun}] = \max\{0.7 \times 0.01, 0.1 \times 0.5\} = 0.05$$

$$m_2[x_2 = \text{rain}] = \max\{0.7 \times 0.01, 0.1 \times 0.8\} = 0.08$$


$$\arg \max_{x_2} m_2[x_2] = \text{rain}$$

$$P(x_1 = \text{sun}, x_2 = \text{rain}, e_1, e_2) = 0.7 \times 0.01 = 0.007$$

best path 1 : t that ends at x_t \neq best path 1 : N that goes through x_t

What do We Do?

$$P(e_1|x_1)P(x_1) \quad P(e_2|x_2)P(x_2|x_1)$$



$$\arg \max_{x_1} m_1[x_1] = ?$$

sun!

$$m_2[x_2] = \max_{x_1} P(e_2|x_2)P(x_2|x_1)m_1[x_1]$$

$$m_2[x_2 = \text{sun}] = \max\{0.7 \times 0.01, 0.1 \times 0.5\} = 0.05$$

$$m_2[x_2 = \text{rain}] = \max\{0.7 \times 0.01, 0.1 \times 0.8\} = 0.08$$

$$\arg \max_{x_2} m_2[x_2] = \text{rain}$$

this path starts at rain


this path starts at rain

idea: what if we also save *where* the best path came from?

$$m_t[x_t] = \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

$$a_t[x_t] = \arg \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

Follow the Breadcrumbs...



for $t = 1$ to N :

$$m_t[x_t] = \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$


$$a_t[x_t] = \arg \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

last state on most likely path: $x_N^* = \arg \max_{x_N} m_N[x_N]$

second to last state on most likely path: $x_{N-1}^* = a_N[x_N^*]$

third to last state on most likely path: $x_{N-2}^* = a_{N-1}[x_{N-1}^*]$

Follow the Breadcrumbs...



for $t = 1$ to N :

$$m_t[x_t] = \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$


$$a_t[x_t] = \arg \max_{x_{t-1}} P(e_t|x_t)P(x_t|x_{t-1})m_{t-1}[x_{t-1}]$$

$$x_N^* = \arg \max_{x_N} m_N[x_N]$$

for $t = N$ to 2:

$$x_{t-1}^* = a_t[x_t^*]$$

Viterbi Algorithm



W_{t-1}	$P(W_t W_{t-1})$	
	sun	rain
sun	0.9	0.1
rain	0.3	0.7


W_t	$P(U_t W_t)$	
	true	false
sun	0.2	0.8
rain	0.9	0.1

Time complexity?
 $O(|X|^2 T)$

Space complexity?
 $O(|X| T)$

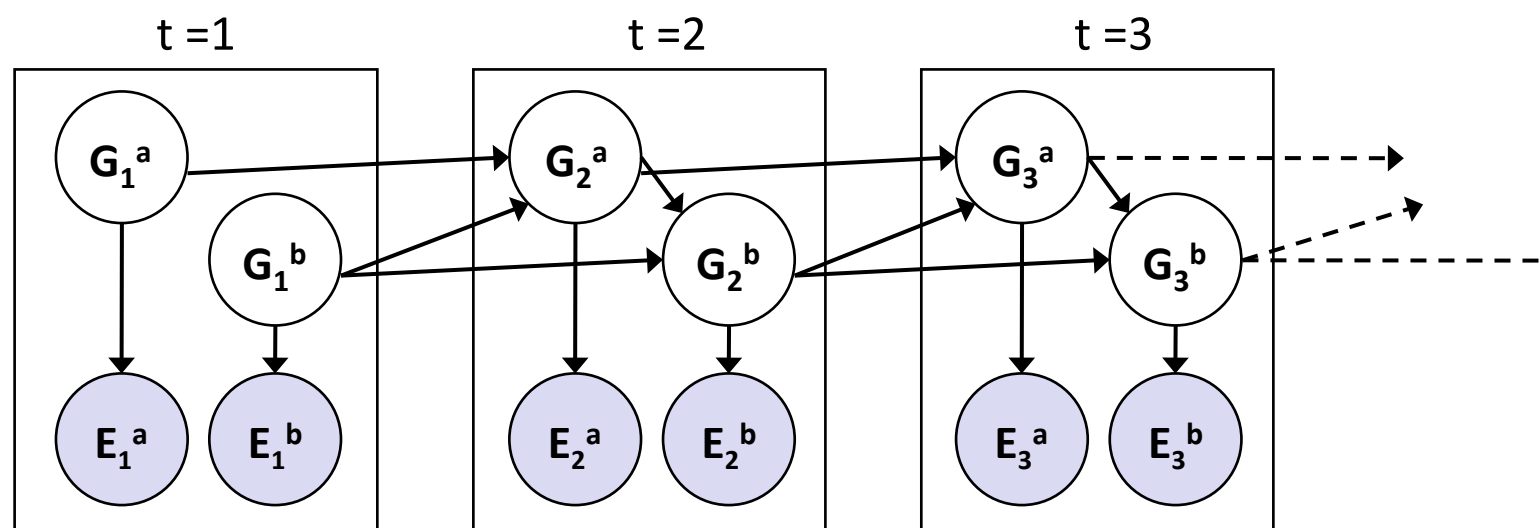

Number of paths?
 $O(|X|^T)$

Dynamic Bayes Nets



Dynamic Bayes Nets (DBNs)


- ❖ We want to track multiple variables over time, using multiple sources of evidence
- ❖ **Idea:** Repeat a fixed Bayes net structure at each time
- ❖ Variables from time t can condition on those from $t-1$



- ❖ Dynamic Bayes nets are a generalization of HMMs

Exact Inference in DBNs

- ❖ Variable elimination applies to dynamic Bayes nets
- ❖ **Procedure:** “unroll” the network for T time steps, then eliminate variables until $P(X_T | e_{1:T})$ is computed



- ❖ **Online belief updates:** Eliminate all variables from the previous time step; store factors for current time only

DBN Particle Filters

- ❖ A particle is a complete sample for a time step
- ❖ **Initialize:** Generate prior samples for the $t=1$ Bayes net
 - ❖ Example particle: $\mathbf{G}_1^a = (3,3)$ $\mathbf{G}_1^b = (5,3)$
- ❖ **Elapse time:** Sample a successor for each particle
 - ❖ Example successor: $\mathbf{G}_2^a = (2,3)$ $\mathbf{G}_2^b = (6,3)$
- ❖ **Observe:** Weight each entire sample by the likelihood of the evidence conditioned on the sample
 - ❖ Likelihood: $P(\mathbf{E}_1^a \mid \mathbf{G}_1^a) * P(\mathbf{E}_1^b \mid \mathbf{G}_1^b)$
- ❖ **Resample:** Select prior samples (tuples of values) in proportion to their likelihood